

Lecture 7: Generative Adversarial Networks

Efstratios Gavves

Lecture overview

- Gentle intro to generative models
- Generative Adversarial Networks
- Variants of Generative Adversarial Networks

Generative models



(a) EBGAN (64x64)



(b) Our results (128x128)

Types of Learning

- Generative modelling
 - Learn the joint pdf: $p(x, y)$
 - Model the world → Perform tasks, e.g. use Bayes rule to classify: $p(y|x)$
 - Naïve Bayes, Variational Autoencoders, GANs
- Discriminative modelling
 - Learn the conditional pdf: $p(y|x)$
 - Task-oriented
 - E.g., Logistic Regression, SVM

Types of Learning


- What to pick?
 - V. Vapnik: “One should solve the [classification] problem directly and never solve a more general [and harder] problem as an intermediate step.”
- Typically, discriminative models are selected to do **the job**
- Generative models give us more theoretical guarantees that the model is going to work as intended
 - Better generalization
 - Less overfitting
 - Better modelling of causal relationships

Applications of generative modeling?

Applications of generative modeling?

- Act as a regularizer in discriminative learning
 - Discriminative learning often too goal-oriented
 - Overfitting to the observations
- Semi-supervised learning
 - Missing data
- Simulating “possible futures” for Reinforcement Learning
- Data-driven generation/sampling/simulation

Applications: Image Generation




(a) Generated by LSGANs.



(b) Generated by DCGANs (Reported in [13]).




Applications: Super-resolution



Applications: Cross-model translation



A map of generative models




Explicit density models

- Plug in the model density function to likelihood
- Then maximize the likelihood


Problems

- Modes must be complex enough → to match data complexity
- Also, model must be computationally tractable
- More details in the next lectures




Generative modeling: Case I

- Density estimation



Train set




Fitted model




Implicit density models

- No explicit probability density function (pdf) needed
- Instead, a sampling mechanism to draw samples from the pdf without knowing the pdf




Implicit density models: GANs

- Sample data in parallel
- Few restrictions on generator model
- No Markov Chains needed
- No variational bounds
- Better qualitative examples
 - Weak but true



Generative modeling: Case II



- Sample Generation



Train examples

Generative modeling: Case II

- Sample Generation



Train examples

New samples (ideally)

What is a GAN?

- **G**enerative

- You can sample novel input samples
- E.g., you can literally “create” images that never existed

- **A**dversarial


- Our generative model G learns adversarially, by fooling an discriminative oracle model D

- **N**etwork


- Implemented typically as a (deep) neural network
- Easy to incorporate new modules
- Easy to learn via backpropagation

GAN: Intuition

- Assume you have two parties
 - Police: wants to recognize fake money as reliably as possible
 - Counterfeiter: wants to make as realistic fake money as possible
- The police forces the counterfeiter to get better (and vice versa)
- Solution relates to Nash equilibrium




GAN: Pipeline



Generator network $x = G(z; \theta^{(G)})$

- Must be differentiable
- No invertibility requirement
- Trainable for any size of z
- Can make conditionally Gaussian given z , but no strict requirement



Generator & Discriminator: Implementation

- The discriminator is just a standard neural network
- The generator looks like an inverse discriminator




Figure 2. Our text-conditional convolutional GAN architecture. Text encoding $\varphi(t)$ is used by both generator and discriminator. It is projected to a lower-dimensions and depth concatenated with image feature maps for further stages of convolutional processing.


[Network Architecture](#)

Training definitions

- Minimax
- Maximin
- Heuristic, non-saturating game
- Max likelihood game

Minimax Game

- $J^{(D)} = -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log D(x) - \frac{1}{2} \mathbb{E}_{z \sim p_z} \log(1 - D(G(z)))$



- $D(x) = 1 \rightarrow$ The discriminator believes that x is a true image
- $D(G(z)) = 1 \rightarrow$ The discriminator believes that $G(z)$ is a true image
- Equilibrium is a saddle point of the discriminator loss
- Resembles Jensen-Shannon divergence
- Generator minimizes the log-probability of the discriminator being correct

[NIPS 2016 Tutorial: Generative Adversarial Networks](#)

A reasonable loss for the generator?

Minimax Game

- For the simple case of zero-sum game

$$J^{(G)} = -J^{(D)}$$

- So, we can summarize game by

$$V(\theta^{(D)}, \theta^{(G)}) = -J^{(D)}(\theta^{(D)}, \theta^{(G)})$$

- Easier theoretical analysis

- In practice not used → when the discriminator starts to recognize fake samples, then ...

Minimax Game

- For the simple case of zero-sum game

$$J^{(G)} = -J^{(D)}$$

- So, we can summarize game by

$$V(\theta^{(D)}, \theta^{(G)}) = -J^{(D)}(\theta^{(D)}, \theta^{(G)})$$

- Easier theoretical analysis


- In practice not used → when the discriminator starts to recognize fake samples, the generator gradients vanish

Heuristic non-saturating game


- $J^{(D)} = -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log D(x) - \frac{1}{2} \mathbb{E}_{z \sim p_z} \log(1 - D(G(z)))$
- $J^{(G)} = -\frac{1}{2} \mathbb{E}_{z \sim p_z} \log(D(G(z)))$

- Equilibrium not any more describable by single loss
- Generator maximizes the log-probability of the discriminator being mistaken
 - Good $G(z)$ \rightarrow $D(G(z)) = 1$ \rightarrow $J^{(G)}$ is maximized
- Heuristically motivated; generator can still learn even when discriminator successfully rejects all generator samples

DCGAN Architecture



Examples



Even vector space arithmetics ...



Man
with
glasses

Man

+



=



Woman with
glasses

Modifying GANs for Max-Likelihood


- $J^{(D)} = -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log D(x) - \frac{1}{2} \mathbb{E}_z \log(1 - D(G(z)))$
- $J^{(G)} = -\frac{1}{2} \mathbb{E}_z \log(\sigma^{-1}(D(G(z))))$
- When discriminator is optimal, the generator gradient matches that of maximum likelihood

[On distinguishability criteria for estimating generative models](#)

Comparison of Generator Losses

*When sample is likely fake, the non-saturating heuristic and the ML cost are flat
→ no gradients in early steps*

*The ML cost variant generates gradients mostly from the “good generations”
→ all gradients from few samples
→ high variance
→ Variance reduction?*




Optimal discriminator

- Optimal $D(x)$ for any $p_{data}(x)$ and $p_{model}(x)$ is always

$$D(x) = \frac{p_{data}(x)}{p_{data}(x) + p_{model}(x)}$$

- Estimating this ratio with supervised learning (discriminator) is the key



Why is this the optimal discriminator?

- $L(D, G) = \int_x p_r(x) \log D(x) + p_g(x) \log(1 - D(x)) dx$
 - Minimize $L(D, G)$ w.r.t. $D \rightarrow \frac{dL}{dD} = 0$ and ignore the integral (we sample over all x)
 - The function $x \rightarrow a \log x + b \log(1 - x)$ attains max in $[0, 1]$ at $\frac{a}{a+b}$
- The optimal discriminator

$$D^*(x) = \frac{p_r(x)}{p_r(x) + p_g(x)}$$

- And at optimality $p_g(x) \rightarrow p_r(x)$, thus

$$\begin{aligned} D^*(x) &= \frac{1}{2} \\ L(G^*, D^*) &= -2 \log 2 \end{aligned}$$

GANs and Jensen-Shannon divergence

- By expanding the Jensen-Shannon divergence, we have

$$\begin{aligned} D_{JS}(p_r || p_g) &= \frac{1}{2} D_{KL}(p_r || \frac{p_r + p_g}{2}) + \frac{1}{2} D_{KL}(p_g || \frac{p_r + p_g}{2}) \\ &= \frac{1}{2} \left(\log 2 + \int_x p_r(x) \log \frac{p_r(x)}{p_r(x) + p_g(x)} dx + \log 2 \right) \end{aligned}$$

GANs and Jensen-Shannon divergence

- By expanding the Jensen-Shannon divergence, we have

$$\begin{aligned} D_{JS}(p_r \parallel p_g) &= \frac{1}{2} D_{KL}(p_r \parallel \frac{p_r + p_g}{2}) + \frac{1}{2} D_{KL}(p_g \parallel \frac{p_r + p_g}{2}) \\ &= \frac{1}{2} \left(\log 2 + \int_x p_r(x) \log \frac{p_r(x)}{p_r(x) + p_g(x)} dx + \log 2 \right) \end{aligned}$$

<https://lilianweng.github.io/lil-log/2017/08/20/from-GAN-to-WGAN.html>

Is the divergence important?

- Does the divergence make a difference?
- Is there a difference between KL-divergence, Jensen-Shannon divergence, ...


$$D_{KL}(p_r || p_g) = \int_x p_r \log \frac{p_r}{p_g} dx$$

$$D_{JS}(p_r || p_g) = \frac{1}{2} D_{KL}(p_r || \frac{p_r + p_g}{2}) + \frac{1}{2} D_{KL}(p_g || \frac{p_r + p_g}{2})$$

- Let's check the KL-divergence

Is the divergence important?

- Forward KL divergence: $D_{KL}(p(x)||q^*(x)) \rightarrow$ high probability everywhere that the data occurs
- Backward KL divergence: $D_{KL}(q^*(x)||p(x)) \rightarrow$ low probability wherever the data does not occur
- Which version makes the model “conservative”?





p_r is what we get and cannot change

p_g is what we make through our model and (through training) change

$$D_{KL}(p_r||p_g) = \int_x p_r \log \frac{p_r}{p_g} dx$$





Is the divergence important?

- $D_{KL}(p(x)||q^*(x)) \rightarrow$ high probability everywhere that the data occurs
- $D_{KL}(q^*(x)||p(x)) \rightarrow$ low probability wherever the data does not occur
- Which version makes the model “conservative”?
- $D_{KL}(q^*(x)||p(x)) = \int q^*(x) \log \frac{q^*(x)}{p(x)}$
 - Avoid areas where $p(x) \rightarrow 0$
- Zero-forcing
 - $q^*(x) \rightarrow 0$ in areas when approximation $\frac{q^*(x)}{p(x)}$ cannot be good




KL vs JS

○ JS is symmetric, KL is not



GAN Problems: Reaching Nash equilibrium causes instabilities

- GANs is a mini-max optimization
 - Non-cooperative game with a tied objective
- Training is not always easy
 - When optimizing one player/network, we might hurt the other one
 - oscillations
- Assume two players $f(x) = xy$
We optimize one step at a time
 - Player 1 minimizes: $\min_x f_1(x) = xy \Rightarrow \frac{df_1}{dx} = y$
 $\Rightarrow x_{t+1} = x_t - \eta \cdot y$
 - Player 2 minimizes: $\min_y f_2(x) = -xy \Rightarrow \frac{df_2}{dx} = -x$
 $\Rightarrow y_{t+1} = y_t + \eta \cdot x$




<https://lilianweng.github.io/lil-log/2017/08/20/from-GAN-to-WGAN.html>

GAN Problems: Vanishing Gradients


$$J^{(D)} = -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log D(x) - \frac{1}{2} \mathbb{E}_z \log(1 - D(G(z)))$$
$$J^{(G)} = -\frac{1}{2} \mathbb{E}_z \log(D(G(z)))$$

- If the discriminator is quite bad
 - no accurate feedback for generator
 - no reasonable generator gradients
- But, if the discriminator is perfect, $D(x) = D^*(x)$
 - gradients go to 0
 - no learning anymore
- Bad when this happens early in the training
 - Easier to train the discriminator than the generator





GAN Problems: Mode collapse

- Very low variability
- It is safer for the generator to produce samples from the mode it knows it approximates well



GAN Problems: Low dimensional supports

- Data lie in low-dim manifolds
- However, the manifold is not known
- During training p_g is not perfect either, especially in the start
- So, the support of p_r and p_g is non-overlapping and disjoint
→ not good for KL/JSD divergences
- Easy to find a discriminating line




Wasserstein GAN

- Instead of KL/JSD, use Wasserstein (Earth Mover's) Distance

$$W(p_r, p_g) = \inf_{\gamma \sim \Pi(p_r, p_g)} \mathbb{E}_{(x,y) \sim \gamma} |x - y|$$

- Even for non-overlapping supports, the distance is meaningful



Feature matching

- Instead of matching image statistics, match feature statistics

$$J^{(D)} = \left\| \mathbb{E}_{x \sim p_r} f(x) - \mathbb{E}_{z \sim p_z} f(G(z)) \right\|_2^2$$

- f can be any statistic of the data, like the mean or the median

Training procedure

- Use SGD-like algorithm of choice
 - Adam Optimizer is a good choice
- Use two mini-batches simultaneously
 - The first mini-batch contains real examples from the training set
 - The second mini-batch contains fake generated examples from the generator
- Optional: run k-steps of one player (e.g. discriminator) for every step of the other player (e.g. generator)

Use labels if possible

- Learning a conditional model $p(y|x)$ is often generates better samples
 - Denton et al., 2015
- Even learning $p(x,y)$ makes samples look more realistic
 - Salimans et al., 2016
- Conditional GANs are a great addition for learning with labels

One-sided label smoothing

- Default discriminator cost:

```
cross_entropy(1., discriminator(data))  
+ cross_entropy(0., discriminator(samples))
```

- One-sided label smoothing:

```
cross_entropy(0.9, discriminator(data))  
+ cross_entropy(0., discriminator(samples))
```

- Do not smooth negative labels:

```
cross_entropy(1.-alpha, discriminator(data))  
+ cross_entropy(beta, discriminator(samples))
```


Benefits of label smoothing

- Max likelihood often is overconfident
 - Might return accurate prediction, but too high probabilities
- Good regularizer
 - Szegedy et al., 2015
- Does not reduce classification accuracy, only confidence
- Specifically for GANs
 - Prevents discriminator from giving very large gradient signals to generator
 - Prevents extrapolating to encourage extreme samples


Batch normalization

- Generally, good practice for neural networks
- Given inputs $X = \{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$
- Compute mean and standard deviation of features of X : μ_{bn}, σ_{bn}
- Normalize features
 - Subtract mean, divide by standard deviation


Batch normalization: Graphically



Batch normalization: Graphically



But, can cause strong intra-batch correlation



Reference batch normalization

- Training with two mini-batches
- One fixed reference mini-batch for computing mean and standard deviation
- The other for doing the training as usual
- Proceed as normal, only use the mean and standard deviation for the batch norm from the fixed reference mini-batch
- Problem: Overfitting to the reference mini-batch

	Standard mini-batch	Reference mini-batch
Iteration 1	$\frac{dJ^{(1)}}{d\theta}$	μ_{bn}, σ_{bn}
Iteration 2	$\frac{dJ^{(2)}}{d\theta}$	μ_{bn}, σ_{bn}
Iteration 3	$\frac{dJ^{(3)}}{d\theta}$	μ_{bn}, σ_{bn}

Solution: Virtual batch normalization

- Mini-batch= standard mini-batch + reference/fixed mini-batch

	Standard mini-batch	Reference mini-batch
Iteration 1	$\frac{dJ^{(1)}}{d\theta}$	$\mu_{bn}^{(R)}, \sigma_{bn}^{(R)}$
Iteration 2	$\frac{dJ^{(2)}}{d\theta}$	$\mu_{bn}^{(R)}, \sigma_{bn}^{(R)}$
Iteration 3	$\frac{dJ^{(3)}}{d\theta}$	$\mu_{bn}^{(R)}, \sigma_{bn}^{(R)}$

Balancing Generator & Discriminator


- Usually the discriminator wins
 - That's good, in that the theoretical justification assume a perfect discriminator
- Usually the discriminator network is bigger than the generator
- Sometimes running discriminator more often than generator works better
 - However, no real consensus
- Do not limit the discriminator to avoid making it too smart
 - Better use non-saturating cost
 - Better use label smoothing

Open Question: Non-convergence

- Optimization is tricky and unstable
 - finding a saddle point does not imply a global minimum
- An equilibrium might not even be reached
- Mode-collapse is the most severe form of non-convergence

Open Question: Mode collapse

- Discriminator converges to the correct distribution
- Generator however places all mass in the most likely point




Open Question: Mode collapse

- Discriminator converges to the correct distribution
- Generator however places all mass in the most likely point
- Problem: low sample diversity



Minibatch features

- Classify each sample by comparing to other examples in the mini-batch
- If samples are too similar, the model is penalized





Open Question: Evaluation of GANs

- Despite the nice images, who cares?
- It would be nice to quantitatively evaluate the model
- For GANs it is even hard to estimate the likelihood

Open Question: Discrete outputs

- The generator must be differentiable
- It cannot be differentiable if outputs are discrete
- E.g., harder to make it work for text
- Possible workarounds
 - REINFORCE [Williams, 1992]
 - Concrete distribution [Maddison et al., 2016]
 - Gumbel softmax [Jang et al., 2016]
 - Train GAN to generate continuous embeddings

Open Question: Semi-supervised classification



Interpretable latent codes

- InfoGAN [Chen et al., 2016]



(a) Azimuth (pose)

(b) Elevation




(c) Lighting

(d) Wide or Narrow

GAN spinoffs

- Conditional GANs
 - Standard GANs have **no encoder!**
- Actor-Critic
 - Related to Reinforcement Learning




Connections to Reinforcement Learning

- GANs interpreted as actor-critic [Pfau and Vinyals, 2016]
- GANs as inverse reinforcement learning [Finn et al., 2016]
- GANs for imitation learning [Ho and Ermin 2016]

Application: Image to Image translation



Application: Style transfer



Application: Face generation

- <https://www.youtube.com/watch?v=XOxxPcy5Gr4>

Summary

- GANs are generative models using supervised learning to approximate an intractable cost function
- GANs can simulate many cost functions, including max likelihood
- Finding Nash equilibria in high-dimensional, continuous, non-convex games is an important open research problem
- GAN research is in its infancy, most works published only in 2016. Not mature enough yet, but very compelling results